



## Written Assignment 2

Yushan Liu   Student ID: 2024214103

November 16, 2024

## Problem 2.1: SVM and Logistic Regression

### (a) Definition of $E_\infty(z)$ and the regularization parameter $\lambda$

The function  $E_\infty(z)$  is defined as:

$$E_\infty(z) = \begin{cases} 0 & \text{if } z \geq 1 \\ 1 - z & \text{if } 0 < z < 1 \\ \infty & \text{if } z \leq 0 \end{cases}$$

Here,  $z$  is the product  $y^{(i)}(\omega^T x^{(i)} + b)$ . This function indicates that:

- If the predicted output is correctly classified with a margin of at least 1 (i.e.,  $z \geq 1$ ), the loss is 0.
- If the predicted output is within the margin (i.e.,  $0 < z < 1$ ), it incurs a linear penalty proportional to how far  $z$  is from 1.
- If the predicted output is incorrect (i.e.,  $z \leq 0$ ), the penalty is infinite, representing a hard constraint violation.

The regularization parameter  $\lambda$  controls the trade-off between maximizing the margin and minimizing the classification error. The constraint for  $\lambda$  is typically:

$$\lambda \geq 0$$

### (b) Definition of $E_{LR}(z)$

Given a target variable  $y \in \{-1, 1\}$ , we know that the logistic regression model defines the probability of  $y = 1$  given the feature vector  $x$  as:

$$p(y = 1|x) = \sigma(w^\top x + b)$$

where  $\sigma(z)$  is the sigmoid function. Consequently the probability of  $y = -1$  is:

$$p(y = -1|x) = 1 - p(y = 1|x) = 1 - \sigma(w^\top x + b) = \sigma(-(w^\top x + b))$$

The likelihood of observing the data for  $m$  samples  $(x^{(i)}, y^{(i)})$  is given by:

$$L(w, b) = \prod_{i=1}^m p(y^{(i)}|x^{(i)})$$

For logistic regression, this can be written as:

$$L(w, b) = \prod_{i=1}^m \sigma(w^\top x^{(i)} + b)^{\frac{1+y^{(i)}}{2}} \cdot \sigma(-(w^\top x^{(i)} + b))^{\frac{1-y^{(i)}}{2}}$$

Thus the negative log-likelihood is:

$$-\log L(w, b) = -\sum_{i=1}^m \left( \frac{1+y^{(i)}}{2} \log \sigma(w^\top x^{(i)} + b) + \frac{1-y^{(i)}}{2} \log \sigma(-(w^\top x^{(i)} + b)) \right)$$

To incorporate the regularization term, we add  $\lambda||w||^2$  to the negative log-likelihood, resulting in the following expression:

$$\sum_{i=1}^m E_{LR}(y^{(i)}(w^\top x^{(i)} + b)) + \lambda||w||^2$$

The function  $E_{LR}(z)$  captures the negative log-likelihood for a single instance and is defined as follows:

$$E_{LR}(z) = \log(1 + e^{-z}) \quad \text{for } z = y^{(i)}(w^\top x^{(i)} + b)$$

In this case:

- If  $y^{(i)} = 1$ ,  $E_{LR}(z) = \log(1 + e^{-(w^\top x^{(i)} + b)})$ .
- If  $y^{(i)} = -1$ ,  $E_{LR}(z) = \log(1 + e^{(w^\top x^{(i)} + b)})$ .

### (c) Definition of $E_{SV}(z)$ and the regularization parameter $\lambda$

The modified SVM optimization problem can be expressed in the desired form, we start by analyzing the introduction of slack variables  $\xi^{(i)}$  in the constraints. The original constraints for hard-margin SVM are replaced by:

$$y^{(i)}(w^\top x^{(i)} + b) \geq 1 - \xi^{(i)}, \quad \text{for } i = 1, \dots, m,$$

where  $\xi^{(i)} \geq 0$  accounts for misclassifications. The new objective function we want to minimize becomes:

$$C \sum_{i=1}^m \xi^{(i)} + \frac{1}{2}||w||^2.$$

To establish a relationship between  $y^{(i)}(w^\top x^{(i)} + b)$  and  $\xi^{(i)}$ , we can write:

$$\xi^{(i)} = \max(0, 1 - y^{(i)}(w^\top x^{(i)} + b)).$$

This formulation shows that:

- If the data point is correctly classified with a margin of at least 1 (i.e.,  $y^{(i)}(w^\top x^{(i)} + b) \geq 1$ ), then  $\xi^{(i)} = 0$ .
- If the data point is on the wrong side of the margin (i.e.,  $y^{(i)}(w^\top x^{(i)} + b) < 1$ ), then  $\xi^{(i)}$  quantifies the extent of misclassification.

The Lagrangian for the soft-margin SVM can be expressed as:

$$\mathcal{L}(w, b, \xi, \alpha) = \frac{1}{2}||w||^2 + C \sum_{i=1}^m \xi^{(i)} - \sum_{i=1}^m \alpha^{(i)} (y^{(i)}(w^\top x^{(i)} + b) - 1 + \xi^{(i)}),$$

where  $\alpha^{(i)}$  are the Lagrange multipliers. Given the relationship established above, we can express the sum of slack variables as:

$$\sum_{i=1}^m \xi^{(i)} = \sum_{i=1}^m E_{SV}(y^{(i)}(w^\top x^{(i)} + b)),$$

where  $E_{SV}(z)$  is defined as follows:

$$E_{SV}(z) = \begin{cases} 0 & \text{if } z \geq 1 \\ 1 - z & \text{if } 0 < z < 1 \\ -\frac{1}{2}z^2 + \frac{1}{2} & \text{if } z < 0 \end{cases}$$

This function captures the penalties incurred based on how far  $y^{(i)}(w^\top x^{(i)} + b)$  deviates from the decision boundary.

The regularization parameter in this context is related to the trade-off between the classification error and the model complexity, and can be expressed as:

$$\lambda = C.$$

## (d) Conclusion and Discussion

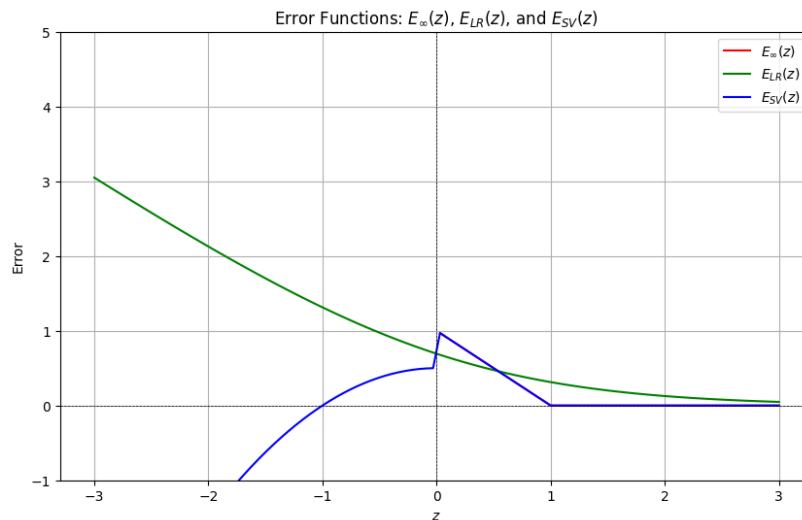
We can use Python libraires like `Numpy` and `Matplotlib` to plot the three error functions above, which implementation code is shown below:

```
def E_inf(z):
    """Error function E∞(·)"""
    return np.where(z >= 1, 0, np.where(z > 0, 1 - z, np.inf))

def E_LR(z):
    """Error function ELR(·)"""
    return np.log(1 + np.exp(-z))

def E_SV(z):
    """Error function ESV(·)"""
    return np.where(z >= 1, 0, np.where(z > 0, 1 - z, -0.5 * z**2 + 0.5))
```

The three error functions are plotted in the figure below:



We can see that:

- $E_\infty(z)$ : It exhibits a hard threshold; any misclassification incurs an infinite penalty. This is ideal for a clear margin but may not be suitable for noisy data.
- $E_{LR}(z)$ : It is smooth and continuous, which allows for a softer transition between correct and incorrect classifications. This smoothness can help the optimization process, leading to better convergence properties.
- $E_{SV}(z)$ : It balances misclassifications and model complexity. The quadratic penalty for  $z < 0$  help avoid extreme penalties while still enforcing the margin constraints.

If the error functions are replaced with other functions, several outcomes may arise:

- **Loss of Robustness:** Using an overly aggressive loss function (like a hard-margin SVM) may lead to poor generalization on noisy or overlapping datasets, resulting in overfitting.
- **Convergence Issues:** Non-smooth functions could lead to difficulties in optimization, causing convergence problems or slow training.
- **Behavior in Overlapping Data:** For functions that do not adequately penalize misclassifications, such as squared error loss in classification tasks, the model may not handle overlap well, potentially leading to high error rates.
- **Interpretability and Application:** Different functions might have different interpretability and applications. For example, soft-margin SVM and logistic regression can be interpreted probabilistically, whereas a hard-margin approach may not.

## Problem 2.2: Naive Bayes Parameter Learning

We are given a dataset  $\{(x^{(i)}, y^{(i)}), i = 1, 2, \dots, m\}$  where each  $x^{(i)}$  is an  $n$ -dimensional vector, with each element  $x_j^{(i)} \in \{0, 1\}$ , and the labels  $y^{(i)} \in \{0, 1\}$ . We have known the following model for Naive Bayes classification:

$y^{(i)}$  follows a Bernoulli distribution:

$$y^{(i)} \sim \text{Bernoulli}(\phi_y)$$

Given  $y^{(i)} = b$ , the features  $x_j^{(i)}$  are independent and follow Bernoulli distributions:

$$x_j^{(i)} | y^{(i)} = b \sim \text{Bernoulli}(\phi_{j|y=b}), \quad b = 0, 1$$

The joint probability distribution of a single data point  $(x^{(i)}, y^{(i)})$  is:

$$p(x^{(i)}, y^{(i)} | \phi_y, \phi_{j|y=b}) = p(y^{(i)} | \phi_y) \cdot \prod_{j=1}^n p(x_j^{(i)} | y^{(i)}, \phi_{j|y=b})$$

Specifically:

- $p(y^{(i)} | \phi_y) = \phi_y^{y^{(i)}} (1 - \phi_y)^{1-y^{(i)}}$
- $p(x_j^{(i)} | y^{(i)} = b, \phi_{j|y=b}) = \phi_{j|y=b}^{x_j^{(i)}} (1 - \phi_{j|y=b})^{1-x_j^{(i)}}$

Thus, the joint probability can be written as:

$$p(x^{(i)}, y^{(i)} | \phi_y, \phi_{j|y=b}) = \phi_y^{y^{(i)}} (1 - \phi_y)^{1-y^{(i)}} \cdot \prod_{j=1}^n \phi_{j|y=y^{(i)}}^{x_j^{(i)}} (1 - \phi_{j|y=y^{(i)}})^{1-x_j^{(i)}}$$

The likelihood function for the entire dataset is the product of the joint probabilities for all data points:

$$L(\phi_y, \phi_{j|y=b}) = \prod_{i=1}^m p(x^{(i)}, y^{(i)} | \phi_y, \phi_{j|y=b}) = \prod_{i=1}^m \left( \phi_y^{y^{(i)}} (1 - \phi_y)^{1-y^{(i)}} \prod_{j=1}^n \phi_{j|y=y^{(i)}}^{x_j^{(i)}} (1 - \phi_{j|y=y^{(i)}})^{1-x_j^{(i)}} \right)$$

To simplify calculations, we take the logarithm of the likelihood function:

$$\begin{aligned} \log L(\phi_y, \phi_{j|y=b}) &= \sum_{i=1}^m (y^{(i)} \log \phi_y + (1 - y^{(i)}) \log(1 - \phi_y)) \\ &\quad + \sum_{i=1}^m \sum_{j=1}^n (x_j^{(i)} \log \phi_{j|y=y^{(i)}} + (1 - x_j^{(i)}) \log(1 - \phi_{j|y=y^{(i)}})) \end{aligned}$$

We now compute the MLE of the parameters by taking the derivative of the log-likelihood function and setting it to zero.

**For  $\phi_y$ :**

$$\frac{\partial \log L}{\partial \phi_y} = \sum_{i=1}^m \left( \frac{y^{(i)}}{\phi_y} - \frac{1 - y^{(i)}}{1 - \phi_y} \right)$$

Setting this to zero, we obtain the MLE for  $\phi_y$ :

$$\hat{\phi}_y = \frac{1}{m} \sum_{i=1}^m y^{(i)}$$

This is the proportion of samples where  $y = 1$ .

**For  $\phi_{j|y=b}$ :**

$$\frac{\partial \log L}{\partial \phi_{j|y=b}} = \sum_{i=1}^m \left( \frac{x_j^{(i)}}{\phi_{j|y=b}} - \frac{1 - x_j^{(i)}}{1 - \phi_{j|y=b}} \right) 1\{y^{(i)} = b\}$$

$$\hat{\phi}_{j|y=b} = \frac{\sum_{i=1}^m 1\{y^{(i)} = b\} x_j^{(i)}}{\sum_{i=1}^m 1\{y^{(i)} = b\}}$$

This is the proportion of times feature  $x_j = 1$  when  $y = b$ .

The maximum likelihood estimates for the parameters are:

$$\hat{\phi}_y = \frac{1}{m} \sum_{i=1}^m y^{(i)}$$

$$\hat{\phi}_{j|y=b} = \frac{\sum_{i=1}^m 1\{y^{(i)} = b\} x_j^{(i)}}{\sum_{i=1}^m 1\{y^{(i)} = b\}}$$

## Problem 2.3: Comparison of Generative and Discriminative Models

### (a) Parameter Estimation for GDA

We are asked to derive the parameter estimation for Gaussian Discriminant Analysis (GDA) with a shared covariance matrix. The setup is as follows:

- $y \sim \text{Bernoulli}(\phi)$
- $x|y = 0 \sim \mathcal{N}(\mu_0, \Sigma)$
- $x|y = 1 \sim \mathcal{N}(\mu_1, \Sigma)$

To estimate the parameters  $\mu_0$ ,  $\mu_1$ ,  $\Sigma$ , and  $\phi$ , we use the maximum likelihood method.

- **Estimate  $\phi$ :** The probability of the label  $y = 1$  is estimated by the fraction of positive labels:

$$\hat{\phi} = \frac{1}{n} \sum_{i=1}^n y^{(i)}$$

- **Estimate  $\mu_0$  and  $\mu_1$ :** The mean vectors  $\mu_0$  and  $\mu_1$  are the empirical means of the samples that have  $y = 0$  and  $y = 1$  respectively:

$$\hat{\mu}_0 = \frac{\sum_{i: y^{(i)}=0} x^{(i)}}{\sum_{i=1}^n 1\{y^{(i)} = 0\}}$$

$$\hat{\mu}_1 = \frac{\sum_{i: y^{(i)}=1} x^{(i)}}{\sum_{i=1}^n 1\{y^{(i)} = 1\}}$$

- **Estimate  $\Sigma$ :** The shared covariance matrix  $\Sigma$  is estimated by pooling the covariance across both classes:

$$\hat{\Sigma} = \frac{1}{n} \sum_{i=1}^n (x^{(i)} - \mu_{y^{(i)}}) (x^{(i)} - \mu_{y^{(i)}})^T$$

where  $\mu_{y^{(i)}}$  is  $\mu_0$  if  $y^{(i)} = 0$ , and  $\mu_1$  if  $y^{(i)} = 1$ .

### (b) Compute for the Parameters

The dataset provided is:

$$\begin{aligned} x^{(1)} &= \begin{pmatrix} 1 \\ 2 \end{pmatrix}, & y^{(1)} &= 0 \\ x^{(2)} &= \begin{pmatrix} 2 \\ 3 \end{pmatrix}, & y^{(2)} &= 0 \\ x^{(3)} &= \begin{pmatrix} 3 \\ 4 \end{pmatrix}, & y^{(3)} &= 1 \\ x^{(4)} &= \begin{pmatrix} 4 \\ 5 \end{pmatrix}, & y^{(4)} &= 1 \end{aligned}$$



- **Estimate  $\phi$ :**

$$\hat{\phi} = \frac{1}{4} (y^{(1)} + y^{(2)} + y^{(3)} + y^{(4)}) = \frac{1}{4} (0 + 0 + 1 + 1) = \frac{1}{2}$$

- **Estimate  $\mu_0$ :**

$$\hat{\mu}_0 = \frac{x^{(1)} + x^{(2)}}{2} = \frac{1}{2} \left( \begin{pmatrix} 1 \\ 2 \end{pmatrix} + \begin{pmatrix} 2 \\ 3 \end{pmatrix} \right) = \begin{pmatrix} 1.5 \\ 2.5 \end{pmatrix}$$

- **Estimate  $\mu_1$ :**

$$\hat{\mu}_1 = \frac{x^{(3)} + x^{(4)}}{2} = \frac{1}{2} \left( \begin{pmatrix} 3 \\ 4 \end{pmatrix} + \begin{pmatrix} 4 \\ 5 \end{pmatrix} \right) = \begin{pmatrix} 3.5 \\ 4.5 \end{pmatrix}$$

- **Estimate  $\Sigma$ :** First, we calculate the covariance terms. For the class  $y = 0$ :

$$(x^{(1)} - \mu_0) = \begin{pmatrix} 1 \\ 2 \end{pmatrix} - \begin{pmatrix} 1.5 \\ 2.5 \end{pmatrix} = \begin{pmatrix} -0.5 \\ -0.5 \end{pmatrix}$$

$$(x^{(2)} - \mu_0) = \begin{pmatrix} 2 \\ 3 \end{pmatrix} - \begin{pmatrix} 1.5 \\ 2.5 \end{pmatrix} = \begin{pmatrix} 0.5 \\ 0.5 \end{pmatrix}$$

For the class  $y = 1$ :

$$(x^{(3)} - \mu_1) = \begin{pmatrix} 3 \\ 4 \end{pmatrix} - \begin{pmatrix} 3.5 \\ 4.5 \end{pmatrix} = \begin{pmatrix} -0.5 \\ -0.5 \end{pmatrix}$$

$$(x^{(4)} - \mu_1) = \begin{pmatrix} 4 \\ 5 \end{pmatrix} - \begin{pmatrix} 3.5 \\ 4.5 \end{pmatrix} = \begin{pmatrix} 0.5 \\ 0.5 \end{pmatrix}$$

Now calculate the covariance matrix:

$$\begin{aligned} \Sigma &= \frac{1}{4} \left( \begin{pmatrix} -0.5 \\ -0.5 \end{pmatrix} \begin{pmatrix} -0.5 & -0.5 \end{pmatrix} + \begin{pmatrix} 0.5 \\ 0.5 \end{pmatrix} \begin{pmatrix} 0.5 & 0.5 \end{pmatrix} + \begin{pmatrix} -0.5 \\ -0.5 \end{pmatrix} \begin{pmatrix} -0.5 & -0.5 \end{pmatrix} + \begin{pmatrix} 0.5 \\ 0.5 \end{pmatrix} \begin{pmatrix} 0.5 & 0.5 \end{pmatrix} \right) \\ &= \begin{pmatrix} 0.25 & 0.25 \\ 0.25 & 0.25 \end{pmatrix} \end{aligned}$$

The computed parameters for this dataset are:

- $\hat{\phi} = 0.5$
- $\hat{\mu}_0 = \begin{pmatrix} 1.5 \\ 2.5 \end{pmatrix}$
- $\hat{\mu}_1 = \begin{pmatrix} 3.5 \\ 4.5 \end{pmatrix}$
- $\hat{\Sigma} = \begin{pmatrix} 0.25 & 0.25 \\ 0.25 & 0.25 \end{pmatrix}$

### (c) Decision Boundary for LDA

The decision boundary is derived by finding the point where the probabilities of the two classes are equal, i.e.,  $P(y = 0|x) = P(y = 1|x)$ . Since both classes follow Gaussian distributions with equal covariance matrices, the decision rule simplifies to:

$$\hat{y} = \hat{y} = \arg \max_y P(y|x) = \arg \max_y \log P(x|y) + \log P(y)$$

This leads to the decision boundary equation:

$$x^T \Sigma^{-1}(\mu_0 - \mu_1) = \frac{1}{2}(\mu_0^T \Sigma^{-1} \mu_0 - \mu_1^T \Sigma^{-1} \mu_1) + \log \left( \frac{\phi}{1 - \phi} \right)$$

### (d) Decision Boundaries Comparion of LDA and LR

#### a. LDA Decision Boundary

From the part (c), we derived that the decision boundary for LDA is given by:

$$x^T \Sigma^{-1}(\mu_0 - \mu_1) = \frac{1}{2}(\mu_0^T \Sigma^{-1} \mu_0 - \mu_1^T \Sigma^{-1} \mu_1) + \log \left( \frac{\phi}{1 - \phi} \right)$$

When  $\Sigma = I$  (Identity Matrix), the decision boundary simplifies to:

$$x^T(\mu_0 - \mu_1) = \frac{1}{2}(\mu_0^T \mu_0 - \mu_1^T \mu_1) + \log \left( \frac{\phi}{1 - \phi} \right)$$

#### b. LR Decision Boundary

Logistic regression models the conditional probability  $P(y = 1|x)$  directly using a sigmoid function:

$$P(y = 1|x) = \frac{1}{1 + \exp(-\theta^T x)}$$

The decision boundary is found by solving:

$$P(y = 1|x) = P(y = 0|x)$$

or equivalently:

$$\theta^T x = 0$$

#### c. Comparison

- Both LDA (with  $\Sigma = I$ ) and LR have **linear decision boundaries**.
- LDA assumes a Gaussian distribution for the data (generative approach), while Logistic Regression does not make any distributional assumptions and models the posterior probabilities directly (discriminative approach).
- If  $\Sigma \neq I$ , the decision boundary for LDA can still be linear but may become more complex depending on the covariance structure, whereas Logistic Regression will always produce a linear boundary.

## (e) LDA vs LR on Small Datasets

### a. LR on Small Datasets

- Logistic Regression is a **discriminative model**, which means it directly models  $P(y|x)$  without making assumptions about the underlying distribution of the data.
- On small datasets, this can be problematic because logistic regression might not have enough data to accurately estimate  $P(y|x)$  without overfitting. It lacks the **regularizing effect** of modeling the data's distribution, which a generative model provides.

### b. LDA on Small Datasets

- LDA is a **generative model**, meaning it models the joint distribution  $P(x, y)$ , specifically  $P(x|y)$  and  $P(y)$ .
- LDA benefits from making stronger assumptions about the data (e.g., Gaussian distribution with shared covariance).

### c. Why Generative Models Might Perform Better

- **With small datasets**, the limited amount of data can make it difficult for discriminative models like logistic regression to accurately learn the conditional probabilities.
- **Generative models**, such as LDA, can perform better because they make use of the **entire data distribution**, imposing structure via assumptions like Gaussianity. This helps stabilize predictions, even with small data.